

Maxwell HPC

Accelerating Machine Learning
- Introductory Guide

Email: aiden.durrant@abdn.ac.uk

Written Guide & Examples: github.com/AidenDurrant/abdn-hpc

What is Maxwell?

Maxwell is a Linux supercomputing cluster housed in the Edward Wright Datacentre.

It is designed for running high computational workloads for simulation, modelling and training.

- NOT a development environment.
- NOT designed to host/interface model inference.

It is funded by University of Aberdeen and National Decommissioning Center and operated by external provider OCF.

Resources Available

CPU (compute) Nodes:

- 11 x Compute Nodes

High Memory Nodes:

- 11 x High Memory Nodes
- 1 x Very High Memory Node (3TB)

GPU Nodes:

- 2 x 2080ti nodes (2 x 12GB RTX 2080ti per node)
- 4 x A100 nodes (3 x 80GB A100 per node)

1 Petabyte of tiered storage

- Each user 1TB

Wide array of HPC optimised software.

- Availability of environment managers

Graphical Interface

- Galaxy Server

SLURM Workload Manager

Slurm is an open source, fault-tolerant, and highly scalable cluster management and job scheduling system for large and small Linux clusters.

Slurms Purpose:

1. Allocation of resources to users for some period of time.
2. Provides a framework for starting, executing, and monitoring work.
3. Manages a queue of pending work.

The short of it:

You interface with the home node, which in-turn communicates with compute nodes to execute your code.

Accessing Maxwell

- Contact Digital Research to gain access to Maxwell
 - Note this is only accessible for PGR, Staff, and under any project where funding is allocated.
- For UG, and PGT the MacLeod server is available.
 - Follow the below link, and fill out the form in the quick links named “Access MacLeod”.

<https://www.abdn.ac.uk/research/digital-research/hpc-for-research-1097.php>

VPN, SSH & SFTP

- The HPC services can be accessed via SSH when connected to the university network (not eduroam).
- Connecting to the university network can be achieved via the f5 VPN if not using a managed device.
 - <https://www.abdn.ac.uk/staffnet/working-here/it-services/remote-access.php>

Maxwell login nodes:

- Maxlogin1.abdn.ac.uk : port 22
- Maxlogin2.abdn.ac.uk : port 22

MacLeod login nodes:

- Macleod1.abdn.ac.uk : port 22
- Macleod2.abdn.ac.uk : port 22

Directory Structure

Two key directories:

- home
- sharedscratch

Home:

Personal 50GB resilient, backed-up.

/uoa/home/<username>/

Sharedscratch:

1TB scratch space for working storage.

/uoa/scratch/users/<username>/

Storage	HPC Backup and Restore Policy
Home Space	<p>Data is backed up as follows:</p> <ul style="list-style-type: none">● Daily backups kept for 14 days<● Weekly backups kept for 1 week <p>Files can be restored as follows:</p> <ul style="list-style-type: none">● To the path that folder/files came, or● To a different specified HPC file path, or● To a specified shared drive <p>Restoration request process:</p> <ul style="list-style-type: none">● Contact servicedesk@abdn.ac.uk
Shared scratch	Not backed up

Data Management

- Transferring of data is best done via SFTP
 - WinSCP on managed devices.
- All code, data, and outputs should be stored in sharedscratch.
 - home should be use only for small quantities of critical data.
- It is recommended to use version control software for code management. (git - github/gitlab)

Maxwell storage is not optimised for speed rather for redundancy.

- Preprocess prior if possible.
- Store data in more efficient formats (i.e. HDF5)
- Avoid multiple jobs accessing same data.

**ALWAYS CHECK WHAT YOU ARE
WRITING TO DISK BEFORE
RUNNING!**

Software

Maxwell is pre-installed with a number of HPC optimized software.

- A list of available software can be found at:
<https://www.abdn.ac.uk/it/documents-uni-only/Maxwell-Galaxy-Software.pdf>
- To utilize software when running a job simply load it via:
`module load <name_of_software>`
- You can load multiple software packages, and unload via:
`module unload <name_of_software>`
- To request any additions to the software list contact the digital research team.

Creating Environments (1/2)

In many cases it can be easier to install your own packages in a virtual environment

For data science workflows that employ Python or R, Anaconda is available.

- Create personal environments.
- Install any python packages and libraries via conda.
- Manage different environments for different projects.

Creating Environments (2/2)

1. `module load miniconda3`
2. `conda create -n <your_env_name> python=3.11`

Note: For first time installs you will be prompted to configure your shell, run the command ``conda init bash`` and then exit / close the terminal, ssh back into maxwell and load the miniconda3 module to continue.

3. `conda activate <your_env_name>`
4. `conda install pytorch torchvision torchaudio cudatoolkit=11.7 -c pytorch`

Note: The current installed version of cuda is 11.7 on maxwell.

Requesting Resources (1/3)

Unlike your standard desktop computer, you instead must submit a “job” for execution to the SLURM scheduler.

Some useful commands are as follows:

- ``sinfo`` - Show summary information
- ``squeue`` - Show job queue
- ``squeue -u <username>`` - Show job queue for a specific user
- ``scontrol show partition a100_full`` - Show gpu partition “a100_full” details
- ``scancel <job_ID>`` - Cancel a job

Requesting Resources (2/3)

To run a job (run your program) on Maxwell you must submit a script to the SLURM scheduler.

The SLURM script must contain 3 things:

1. Define the resource requirements for the job.
2. Activate the environment we created earlier.
3. Specify the script we wish to run.

```
1.  #!/bin/bash
2.  #SBATCH --nodes=1 # number of nodes
3.  #SBATCH --cpus-per-task=12 # number of cores
4.  #SBATCH --mem=32G # memory pool for all cores

5.  #SBATCH --ntasks-per-node=1 # one job per node
6.  #SBATCH --gres=gpu:2 # two gpus per node
7.  #SBATCH --partition=a100_mig

8.  #SBATCH -o slurm.%j.out # STDOUT
9.  #SBATCH -e slurm.%j.err # STDERR

10. #SBATCH --mail-type=ALL
11. #SBATCH --mail-user=<username>@abdn.ac.uk

12. module load anaconda3
13. source activate test

14. srun python example_script.py --epochs=25 --save /home/<username>/sharedscratch/
```

Requesting Resources (3/3)

- The partition specifies which compute nodes you will use, e.g. hmem, cgpu, a100, etc.
- To submit the previous job script named `run_example.sh` to the Slurm scheduler we use the `sbatch` command:

```
`sbatch run_example.sh`
```

- The job may run immediately or may take upto a few days to run depending on the number of existing jobs, you will receive an email notification if you use the SLURM command `email`.
- To check the status of queued and running jobs, use the following:

```
`squeue -u <username>`
```

- If you wish to cancel a job simply use:

```
`scancel <job_ID>`
```

Other SLURM Settings

SLURM arrays:

A slurm array batch job is similar to just running a 'for' loop over the sbatch script.

- All runs have the same jobid with a predictable id as a suffix.
- #SBATCH --array=1-10

```
srun python addone.py $SLURM_ARRAY_TASK_ID
```

#SBATCH Commands:

A full list can be found here: <https://slurm.schedmd.com/sbatch.html>

Monitoring Jobs

Monitoring jobs can most simply be achieved by printing the logs specified in the run script.

```
#SBATCH -o slurm.%j.out  
#SBATCH -e slurm.%j.err
```

This produces two logs per submitted job:

- STDOUT
- STDERR

Many other web/cloud monitoring tools exist (ML focused):

- Weights and Biases - <https://wandb.ai>
- Neptune.ai - <https://neptune.ai>

Wall Clock Time

The duration for which the nodes remain allocated to you.

- It is important to note that all jobs have a maximum walltime of 24 hours.

```
#SBATCH -t 24:00:00
```

If your job requires computation that is longer than 24 hours, you must handle it within the code/script.

- Checkpointing
 - Save and Resume
- Detection of resume checkpoint
- Auto-Requeue SLURM

```
1. timeout 23h ./the_program
2. if [[ $? == 124 ]]; then
3.   scontrol requeue $SLURM_JOB_ID
4. fi
```

Distributed Computation

Using more than 1 GPU

- When requesting more than 1 GPU ensure your code is configured to utilize the additional GPU's
- By default PyTorch will only use 1 GPU.
- To use multiple GPUs on one node implement DataParallel.

<https://pytorch.org/docs/stable/generated/torch.nn.DataParallel.html>

Using more than 1 Node (*Advanced*)

- When requesting more than 1 node ensure your code is written with distributed computation.
- By default PyTorch will run separate processes on each node.
- To use multiple nodes, ensure your code communicates across nodes.

https://pytorch.org/tutorials/intermediate/ddp_tutorial.html

Additional Resources

- SLURM: <https://slurm.schedmd.com/quickstart.html>
- Interactive Allocation: <https://engaging-web.mit.edu/eofe-wiki/slurm/srun/>
- Jupyter Notebooks: <https://nero-docs.stanford.edu/jupyter-slurm.html>
- ChatGPT for SLURM
- University Information: <https://www.abdn.ac.uk/research/digital-research/hpc-for-research-1097.php>
- Written Guide: <https://github.com/AidenDurrant/abdn-hpc>